

# Midterm Notes

## What is a Database Management System (DBMS)?

- Massive
  - A large amount of data is being managed (i.e. 2 terabytes)
  - Data cannot be stored in main memory, stored on physical disks somewhere
- Persistent
  - Data need to be available in a consistent state
- Convenient
  - SQL allows for dynamic on-the-fly searches
- Efficient
  - The system understands the language and comes up with the right way of executing queries efficiently (query optimization)
  - Standard benchmark (TPC) that estimates how fast the database system runs
- Safe
  - Safety from system failure
    - ❖ If a web server crashes, or even if the database server crashes, the data is still safe and unharmed
  - Safety from malicious users
    - ❖ The system shows the right sets of the data to the right users
- Multi-User
  - Example: A husband and wife have \$300 in their bank account. John withdraws \$100 and Susan withdraws \$50 at different locations at the same time. What is the remaining balance left in the account?
- Location of data – In concurrent programming, the data is in memory, where as in DBMS, the data is in a physical disk
- Granularity of data – In a file system, the granularity of data is a file/page, in DBMS, the granularity of data is in tuples

## Database construction steps

- Domain analysis
- Database design
  - Entity-Relational model – Analyze the relations in the database
- Table creation
  - Data Definition Language (DDL)
- Data load
  - No standard SQL command to load bulk data
- Query & update
  - Data Manipulation Language (DML)

## Relational algebra

- High level concept
  - Relational model stores tables into the disk
  - Generates relations based on the query that is executed upon the disk
- Query – a question
  - Poses a query to the underline system, and the system returns an answer
- Operators
  - Selection operator “ $\sigma$ ”: Returns the tuples in the table where the tuples meets the requirements specified by the operator
    - ❖ i.e., Find all students that are younger than 18  $\equiv \sigma_{\text{age} < 18}(\text{Student})$
    - ❖ i.e., Find all students that are younger than 18 and has a GPA greater than 3.7  $\equiv \sigma_{\text{age} < 18 \wedge \text{GPA} > 3.7}(\text{Student})$
    - ❖ Valid Boolean operators:  $>$ ,  $<$ ,  $=$ ,  $\neq$ ,  $\geq$ ,  $\leq$ ,  $\wedge$ ,  $\vee$ ,  $\neg$
  - Projection operator “ $\pi$ ”: Returns the chosen attributes in a table
    - ❖ i.e., Find the department of all courses  $\equiv \pi_{\text{dept}}(\text{Course})$
    - ❖ i.e., Find the name, address, and GPA of all students  $\equiv \pi_{\text{name, addr, GPA}}(\text{Student})$
  - Cross-product operator “ $\times$ ”: Returns all possible combinations of the rows in two tables
    - ❖ i.e., Find all combinations of students and enrolled students  $\equiv \text{Student} \times \text{Enroll} = (\text{Student.sid, name, addr, age, GPA, Enroll.sid, dept, cnum, sec})$
    - ❖ Number of resulting tuples is  $|\text{table}_1| * |\text{table}_2|$  (i.e.,  $4 * 4 = 16$ )
    - ❖ Conflicts in attribute names are prefaced by the table name
  - (Natural) Join operator “ $\bowtie$ ”: Returns the cross-products of table tables where common attributes match in value

- ❖ i.e., Find all the enrolled students and the course that they are enrolled in  $\equiv \text{Enroll} \bowtie \text{Course} \equiv \sigma_{\text{Enroll.dept} = \text{Course.dept} \wedge \text{Enroll.cnum} = \text{Course.cnum} \wedge \text{Enroll.sec} = \text{Course.sec}}(\text{Enroll} \times \text{Course})$
- $\theta$ -join operator “ $\bowtie_{\theta}$ ”: Returns tuples in a cross-product that meets the selection criteria
  - ❖ i.e., Find all combinations of students with a GPA higher than 3.7 and enrolled students  $\equiv \text{Student} \bowtie_{\text{Student.sid} = \text{Enroll.sid} \wedge \text{GPA} > 3.7} \text{Enroll} \equiv \sigma_{\text{Student.sid} = \text{Enroll.sid} \wedge \text{GPA} > 3.7}(\text{Student} \times \text{Enroll})$
- Set union operator “ $\cup$ ”: Returns the union of two tables
  - ❖  $\text{Student}(\text{sid, name, addr, age, GPA})$
  - ❖  $\text{GradStudent}(\text{sid, name, addr, age, GPA, advisor})$
  - ❖ i.e., Find the SIDs of all the students  $\equiv \pi_{\text{sid}}(\text{Student}) \cup \pi_{\text{sid}}(\text{GradStudent})$
- Set difference operator “ $-$ ”: Return the difference of two tables
  - ❖ i.e., Find the department, course number, and section of all classes without enrolled students  $\equiv \pi_{\text{dept, cnum, sec}}(\text{Class}) - \pi_{\text{dept, cnum, sec}}(\text{Enroll})$
  - ❖ i.e., Find the titles of all the classes without enrolled students  $\equiv \pi_{\text{title}}((\pi_{\text{dept, cnum, sec}}(\text{Class}) - \pi_{\text{dept, cnum, sec}}(\text{Enroll})) \bowtie \text{Class})$
- Set intercept operator “ $\cap$ ”: Returns the interception of two tables
  - ❖ i.e., Find all the instructors who teaches both CS and EE classes  $\equiv \pi_{\text{instructor}}(\sigma_{\text{dept}='CS'}(\text{Class})) \cap \pi_{\text{instructor}}(\sigma_{\text{dept}='EE'}(\text{Class}))$
  - ❖  $R \cap S = R - (R - S)$
- Rename operator “ $\rho$ ”: Renames a table to another name
  - ❖ i.e., Find pairs of student names who live at the same address  $\equiv \sigma_{R.\text{name} > S.\text{name}}(\pi_{\text{name, name}}(\sigma_{R.\text{addr} = S.\text{addr}}(\rho_R(\text{Student}) \times \rho_S(\text{Student}))))$
- Division operator “ $/$ ”: Returns tuples that exist in all of another tuple/table
  - ❖ i.e., Find the SIDs of students who is enrolled in all of the CS classes
    - Things we’d need to know:
      - ✓ All possible combinations of students taking CS classes  $\equiv \pi_{\text{sid}}(\text{Student} \times \pi_{\text{dept, cnum, sec}}(\sigma_{\text{dept}='CS'}(\text{Class}))) - \text{Enroll}$
      - ✓ Students who are not taking at least one of the classes  $\equiv$  All possible combinations – enrolled students
      - ✓ Students who are taking all CS classes  $\equiv$  All students – students who are not taking at least one of the classes
    - $\pi_{\text{sid}}(\text{Student}) - \pi_{\text{sid}}(\pi_{\text{sid}}(\text{Student} \times \pi_{\text{dept, cnum, sec}}(\sigma_{\text{dept}='CS'}(\text{Class}))) - \text{Enroll})$
  - ❖  $R/S = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$
- Common mistake: Questions that need to use the logical opposite approach
  - ❖ i.e., Find the SIDs of students who are not taking any CS classes
    - $\pi_{\text{sid}}(\sigma_{\text{dept} \neq 'CS'}(\text{Student} \bowtie \text{Enroll})) \dots$  WRONG!
    - $\pi_{\text{sid}}(\text{Student}) - \pi_{\text{sid}}(\sigma_{\text{dept}='CS'}(\text{Enroll})) \dots$  CORRECT!

## Standard Query Language (SQL)

- Operators (multi-set or bag semantic):
  - SELECT attribute: Specify the chosen attribute in a table
  - FROM table: Specify the table to execute the query
  - WHERE condition: Specify the condition to choose tuples
  - ORDER BY attribute [ASC/DESC]: Specify an attribute to order the result by
  - DISTINCT: Eliminate duplicate tuples in result
  - Wildcards “ $\_$ ” and “ $\%$ ” – “ $\_$ ” denotes single character, “ $\%$ ” denotes any number of characters
  - LIKE: Matches substrings
  - GROUP BY: Specify an attribute to group by
  - HAVING: Specify a condition on aggregates to choose by
- Set operators (set semantic):
  - UNION: Returns a table that is the first table and the second table appended together
  - INTERSECT: Returns the tuples in the first table that also exist in the second table
  - DIFFERENCE: Returns the tuples in the first table subtract the second table

- EXCEPT: Returns the tuples in the first table that is not in the second table
- NOT IN: Returns true if tuple is not in a specified set
- ALL: Returns true if tuple compares true to all elements of a set
- SOME: Returns true if tuple compares true to at least one element of a set
- Subqueries
  - i.e., Give me the names of students who are not taking any CS classes

```
SELECT name
FROM Student
WHERE sid NOT IN
  (SELECT sid
   FROM Enroll
   WHERE dept='CS')
```

- Aggregation operators
  - AVG: Returns the average of the selected attributes, does not average null values
  - COUNT: Returns the count of the selected attributes, counts null values
  - MAX/MIN: Returns the max or min of the selected attributes, does not consider null values
- Database modification operators
  - INSERT INTO table VALUES (attr<sub>1</sub>, ...)
  - INSERT INTO table query
  - UPDATE table SET attr<sub>i</sub> = new\_value [WHERE condition]
  - DELETE FROM table WHERE condition

### Integrity constraints

- Key constraints
  - Defines the keys of each table, which uniquely identifies a tuple
  - If a key is not defined as NOT NULL, then DB2 will give an error prompting the key to be defined as NOT NULL
  - i.e., Course(dept, cnum, sec, unit, instructor, title)

```
CREATE TABLE Course (
  dept CHAR(2) NOT NULL,
  cnum INT NOT NULL,
  sec INT NOT NULL,
  unit INT,
  instructor VARCHAR(30),
  title VARCHAR(30),
  PRIMARY KEY(dept, cnum, sec),
  UNIQUE(dept, cnum, instructor),
  UNIQUE(dept, sec, title) )
```

- Referential integrity
  - Values of one attribute values in one table will always appear in the values of an attribute of another table (i.e., Enroll.sid will always appear in Student.sid)
  - Value of a foreign key can be NULL
  - Operators that may cause violations
    - DELETE from referenced table
    - UPDATE on referenced table
    - INSERT into referencing table
    - UPDATE on referencing table
  - When changing the referencing table, the DBMS may abort the command and not allow the change
  - i.e., sid must be an entry in Student; dept, cnum, sec must be entries in Class

```
CREATE TABLE Enroll (
  sid INT REFERENCES Student(sid),
  dept CHAR(2),
  cnum INT,
  sec INT,
  FOREIGN KEY (dept, cnum, sec)
    REFERENCES Class(dept, cnum, sec) |
    ON DELETE CASCADE
    ON UPDATE SET NULL )
```

- Referencing attributes called FOREIGN KEY
- Referenced attributes must be PRIMARY KEY or UNIQUE
- Tuple (CHECK) constraint
  - Enforced for every tuple within a particular table
  - CHECK(<condition>) in CREATE TABLE
    - i.e., The units of all CS classes are above 3

```
CRATE TABLE Student(
  dept CHAR(2),
  cnum INT,
  unit INT,
```

```
title VARCHAR(50),
CHECK (dept <> 'CS' OR unit > 3) )
```

- Constraint is checked whenever a tuple is inserted/updated.
- In SQL92, conditions can be more complex, e.g., with subqueries
- SQL assertions
  - Constraint on the entire relation or database
  - CREATE ASSERTION <assertion\_name> CHECK (<condition>)
  - i.e., Average GPA > 3.0

```
CREATE ASSERTION HighGPA CHECK (
  3.0 < (SELECT AVG(GPA) FROM Student) )
```

### Disks

- A typical disk
  - Platter diameter: 1-5 in
  - Cylinders: 100 – 2000
  - Platters: 1 – 20
  - Sectors per track: 200 – 500
  - Sector size: 512 – 50K
  - Overall capacity: 1G – 200GB
    - (sectors / track) × (sector size) × (cylinders) × (2 × number of platters)
- Disk access time
  - Access time = (seek time) + (rotational delay) + (transfer time)
    - Seek time – moving the head to the right track
    - Rotational delay – wait until the right sector comes below the head
    - Transfer time – read/transfer the data
- Seek time
  - Time to move a disk head between tracks
    - Track to track ~ 1ms
    - Average ~ 10 ms
    - Full stroke ~ 20 ms
- Rotational delay
  - Typical disk:
    - 3600 rpm – 15000 rpm
    - Average rotational delay?
      - 3600 rpm / 60 sec = 60 rps; delay = 1/60 sec
- Transfer rate
  - Burst rate
    - (# of bytes per track) / (time to rotate once)
  - Sustained rate
    - Average rate that it takes to transfer the data
    - (# of bytes per track) / (time to rotate once + track-to-track seek time)
- Abstraction by OS
  - Sequential blocks – No need to worry about head, cylinder, sector
  - Access to random blocks – Random I/O
  - Access to consecutive blocks – Sequential I/O
- Random I/O vs. Sequential I/O
  - Assume
    - 10ms seek time
    - 5ms rotational delay
    - 10MB/s transfer rate
    - Access time = (seek time) + (rotational delay) + (transfer time)
  - Random I/O
    - Execute a 2K program – Consisting of 4 random files (512 each)
    - (( 10ms ) + ( 5ms ) + ( 512B / 10MB/s )) × 4 files = 60ms
  - Sequential I/O
    - Execute a 200K program – Consisting of a single file
    - ( 10ms ) + ( 5ms ) + ( 200K / 10MB/s ) = 35ms
- Block modification
  - Byte-level modification not allowed
    - Can be modified by blocks
  - Block modification
    - Read the block from disk
    - 2. Modify in memory
    - 3. Write the block to disk
- Buffer, buffer pool
  - Keep disk blocks in main memory
    - Avoid future read
    - Hide disk latency
  - Buffer, buffer pool
    - Dedicated main memory space to “cache” disk blocks
    - Most DBMS let users change buffer pool size

## Files

- Spanned vs. Unspanned
  - Unspanned – Store as many tuples into a block, forget about the extra remaining space
  - Spanned – Store as many tuples into a block, store part of the next tuple into the block
- Deletion
  - For now, ignore spanning issue, irrelevant for current discussion
  - What should we do?
    - ❖ Copy the last entry into the space
    - ❖ Shift all entries forward to fill the space
    - ❖ Leave it open and fill it with the next update
      - Have a pointer to point to the first available empty slot
      - Have a bit-map of the occupancy of the tuples
- Variable-Length Tuples
  - Reserved Space – Reserve the maximum space for each tuple
  - Variable Length
    - ❖ Tuple length in the beginning
    - ❖ End-of-record symbol
    - ❖ Pack the tuples tightly into a page
  - Update on Variable Length Tuples?
    - ❖ If new tuple is shorter than the tuple before – just place it where it was
    - ❖ If new tuple is longer than the tuple before – delete the old tuple and place it at the end of the block with free space
  - Slotted Page
    - ❖ Header slots in the beginning, pointing to tuples stored at the end of the block
- Long Tuples
  - Spanning
  - Splitting tuples – Split the attributes of tuples into different blocks
- Sequential File – Tuples are ordered by some attributes (search key)
- Sequencing Tuples
  - Inserting a new tuple
    - ❖ Easy case – One tuple has been deleted in the middle
      - Insert new tuple into the block
    - ❖ Difficult case – The block is completely full
      - May shift some tuples into the next block, if there are space in the next block
      - If there are no space in the next block, use the overflow page
  - Overflow page
    - ❖ Overflow page may overflow as well
      - Use pointers to point to additional overflow pages
      - May slow down performance, because this uses random access
  - Any problem?
    - ❖ PCTFREE in DBMS
      - Keeps a percentage of free space in blocks, to reduce the number of overflow pages
      - Not a SQL standard
- Example:
  - ❖ 1,000,000 records (900-bytes/rec)
  - ❖ 4-byte search key, 4-byte pointer
  - ❖ 4096-byte block
- How many blocks for table?
  - ❖  $\text{Tuples / block} = \text{size of block} / \text{size of tuples} = 4096 / 900 = 4 \text{ tuples}$
  - ❖  $\text{Records / tuples} = 1,000,000 / 4 = 250,000 \text{ blocks}$
  - ❖  $250,000 \text{ blocks} * 4096 \text{ bytes / block} = 1\text{GB}$
- How many blocks for index?
  - ❖  $\text{Index / block} = 4096 / 8 = 512$
  - ❖  $\text{Records / indexes} = 1,000,000 / 512 = 1956$
  - ❖  $1956 \text{ blocks} * 4096 \text{ bytes / block} = 8\text{MB}$
- Sparse index
  - For every block, create an index entry which to search on, and a pointer to the block that it points to (even smaller index size of the dense index)
  - In real world, this reduces the index size dramatically, because there may be many tuples in one block, for which sparse index only creates one index entry to those tuples
- Sparse 2<sup>nd</sup> level
  - For every index block, create an index entry which to search on, and a pointer to the index block that it points to (an index on the index, which further reduces in size)
  - Can create multiple level of indexes (multi-level index)
- Terms
  - Index sequential file (Index Sequential Access Method)
  - Search key ( $\neq$  primary key)
  - Dense index vs. Sparse index
  - Multi-level index
- Duplicate keys
  - Dense index, one way to implement – Create an index entry for each tuple
  - Dense index, typical way – Create an index entry for each unique tuple
- Updates on the index?
  - Insertion (empty) – First follow the link structure to identify where the tuple should be located (found with enough space)
  - Insertion (overflow) – Create an overflow block with a pointer from the original block, which adds the entry 15 into the overflow block
  - Insertion (redistribute)
    - ❖ Try to move blocks to other adjacent blocks
    - ❖ Update any changes to the indexes as needed
- Deletion (one tuple)
  - See which index block the tuple is located
  - If the first entry of the block is not deleted, no update to index necessary
  - If the first entry of the block is deleted, update the index appropriately
- Deletion (entry block)
  - If the entire block is deleted, the index entry can be deleted
  - Move all index entries within the block up to compact space
- Primary index – Index that is created over the set of attributes that the table is stored (also called clustering index)
- Secondary index
  - Index on a non-search-key
  - Unordered tuples – Non-sequential file
  - Sparse index make sense?
    - ❖ Does not make sense because the files are not in sequence
  - Dense index on first level
  - Sparse index from the second level
- Duplicate values & secondary indexes
  - One option – Dense index for every tuple that exist
  - Buckets
    - ❖ Blocks that holds pointers to the same index keys
    - ❖ Intermediary level between the index and the tables
- Traditional index
  - Advantage
    - ❖ Simple
    - ❖ Sequential blocks
  - Disadvantage
    - ❖ Not suitable for updates
    - ❖ Becomes ugly (loses sequentiality and balance) over time

## Indexing

- Basic idea – Build an “index” on the table
    - An auxiliary structure to help us locate a record given to a “key”
    - Example: User has a key (40), and looks up the information in the table with the key
  - Indexes to learn
    - Tree-based index
      - ❖ Index sequential file
        - Dense index vs. sparse index
        - Primary index (clustering index) vs. Secondary index (non-clustering index)
      - ❖ B+ tree
    - Hash table
      - ❖ Static hashing
      - ❖ Extensible hashing
  - Dense index
    - For every tuple in the table, create an index entry which to search on, and a pointer to the tuple that it points to (so just an index with pointers to the tuple in the block that the tuple is in)
    - Dense index blocks contain more indexes per block than tuples in their blocks, because dense indexes are much smaller in size than the tuple that they point to.
  - Why dense index?
- B+ tree**
- B+ tree
    - Most popular index structure in RDBMS
    - Advantage

- ❖ Suitable for dynamic updates
- ❖ Balanced
- ❖ Minimum space usage guarantee
- Disadvantage
  - ❖ Non-sequential index blocks
- B+ tree example
  - N pointers, (n-1) keys per node
  - Keys are sorted within a node
  - Balanced: all leaves at same level
- Same non-leaf node (n = 3)
  - At least  $\lceil n/2 \rceil$  pointers (except root)
  - At least 2 pointers in the root
- Nodes are never too empty
  - Use at least
    - ❖ Non-leaf:  $\lceil n/2 \rceil$  pointers
    - ❖ Leaf:  $\lceil (n-1)/2 \rceil + 1$  pointers
- Insert into B+ tree (simple case)
- Insert into B+ tree (leaf overflow)
  - Split the leaf, insert the first key of the new node
  - Move the second half to a new node
  - Insert the first key of the new node to the parent
- Insert into B+ tree (non-leaf overflow)
  - Find the middle key
  - Move everything on the right to a new node
  - Insert (the middle key, the pointer to the new node) into the parent
- Insert into B+ tree (new root node)
  - Insert (the middle key, the pointer to the new node) into the new root
- Delete from B+ tree (simple case)
  - Underflow (n = 4)
    - ❖ Non-leaf  $< \lceil n/2 \rceil = 2$  pointers
    - ❖ Leaf  $< \lceil (n-1)/2 \rceil + 1 = 3$  pointers
- Delete from B+ tree (coalesce with sibling)
  - Move the node across to its sibling if there are rooms available
- Delete from B+ tree (re-distribute)
  - Grab a key from the sibling and move it to the underflowing node
- Delete from B+ tree (coalesce at non-leaf)
  - Push down the parent key into the child node
  - Get the mid-key from parent
  - Push down one of the grand-parent keys into the neighboring parent key
  - Point the child key to the push-down grand-parent key
- Delete from B+ tree (redistribute at non-leaf)
  - Combine the parent and the neighboring keys to make one full node
  - Push down one of the grand-parent keys
  - Push up one of the neighboring parent keys
- B+ tree deletions in practice
  - Coalescing is often not implemented
    - ❖ Too hard and not worth it!
- Question on B+ tree
  - SELECT \*  
FROM Student  
WHERE sid > 60
  - Very efficient on B+ tree
  - Not efficient with hash tables
- Index creation in SQL
  - CREATE INDEX ON <table> (<attr>, <attr>, ...)
  - i.e.,
 

```
CREATE INDEX ON
Student (sid)
```

    - ❖ Creates a B+ tree on the attributes
    - ❖ Speeds up lookup on sid
  - Clustering index (in DB2)
 

```
CREATE INDEX ON
Student (sid) CLUSTER
```

    - ❖ Tuples are sequenced by sid

## Hash table

- What is a hash table?
  - Hash table

- ❖ Hash function
  - Divide the integer by the key
  - $h(k)$ : key  $\rightarrow$  integer  $[0 \dots n]$
  - i.e.,  $h(\text{'Susan'}) = 7$
- ❖ Array for keys:  $T[0 \dots n]$
- ❖ Given a key  $k$ , store it in  $T[h(k)]$
- Properties
  - ❖ Uniformity – entries are distributed across the table uniformly
  - ❖ Randomness – even if two keys are very similar, the hash values will eventually be different
- Why hash table?
  - Direct access
  - saved space – Do not reserve a space for every possible key
- Hashing for DBMS (static hashing)
  - Search key  $\rightarrow h(\text{key})$ , which points to a (key, record) in disk blocks (buckets)
- Record storage
  - Can store as whole record or store as key and pointer, which points to the record
- Overflow
  - Size of the table is fixed, thus there is always a chance that a bucket would overflow
  - Solutions:
    - ❖ Overflow buckets (overflow block chaining) – link to an additional overflow bucket
      - More widely used
    - ❖ Open probing – go to the next bucket and look for space
      - Not used very often anymore
- How many empty slots to keep?
  - 50% to 80% of the blocks occupied
    - ❖ If less than 50% used
      - Waste space
      - Extra disk look-up time with more blocks that needs to be looked up
    - ❖ If more than 80% used
      - Overflow likely to occur
- Major problem of static hashing
  - How to cope with growth?
    - ❖ Data tends to grow in size
    - ❖ Overflow blocks unavoidable
- Extensible hashing
  - Two ideas
    - ❖ Use  $i$  of  $b$  bits output by hash function
      - Use the prefix of the first  $i$  bits of a string of  $b$ -bits in length (i.e., use the first 3 bits of a 5-bit hash value)
    - ❖ Use directory that maintains pointers to hash buckets (indirection)
      - Maintain a directory and do some indirection
  - Possible problems
    - ❖ When there are many duplicates, because there are more copies than digits (?)
      - Still need to use overflow buckets
    - ❖ No space occupancy guarantee when values are extremely skewed, thus needing a very good hash function
    - ❖ Efficient for equality operator (=), but not efficient for range operators (>, <, etc.)
  - Bucket merge
    - ❖ Bucket merge condition
      - Bucket  $i$ 's are the same
      - First  $(i-1)$  bits of the hash key are the same
    - ❖ Directory shrink condition
      - All bucket  $i$ 's are smaller than the directory  $i$
  - Summary
    - ❖ Can handle growing files
      - No periodic reorganizations
    - ❖ • Indirection
      - Up to 2 disk accesses to access a key
    - ❖ • Directory doubles in size
      - Not too bad if the data is not too large